

Need for Speed

NumPy@CAID


Outline

- Introduction What is CAID?
- Problem statement It's too slow!
- Solution NumPy
- Results Now it's fast

Introduction

What is CAID?


CAID



Amino Acids

- True = DisProt
- Pred = published predictors

CAID




Amino Acids

- True = DisProt
- Pred = published predictors
- Threshold


→
true: 0001110111
pred: 1100000011

CAID




Amino Acids

- True = DisProt
- Pred = published predictors
- Threshold



CAID



Amino Acids

- True = DisProt
- Pred = published predictors
- Threshold

true: 0001110111
pred: 1100000011

	0	1
0	TN=2	FP=2
1	FN=4	TP=2

Confusion matrix

$$ACC = \frac{TP + TN}{TP + TN + FP + FN} = \frac{2 + 2}{2 + 2 + 4 + 2} = 0.40$$

Metrics (e.g. accuracy)

Problem Statement

It's too slow!

Code complexity

Dimension	Elements in dimension	Cumulative product	Order of magnitude	Execution time
Metrics	10	10	10^1	~1 second

Code complexity

Dimension	Elements in dimension	Cumulative product	Order of magnitude	Execution time
Metrics	10	10	10^1	~1 second
Proteins	646	~6,500	10^3	~1 minutes

Code complexity

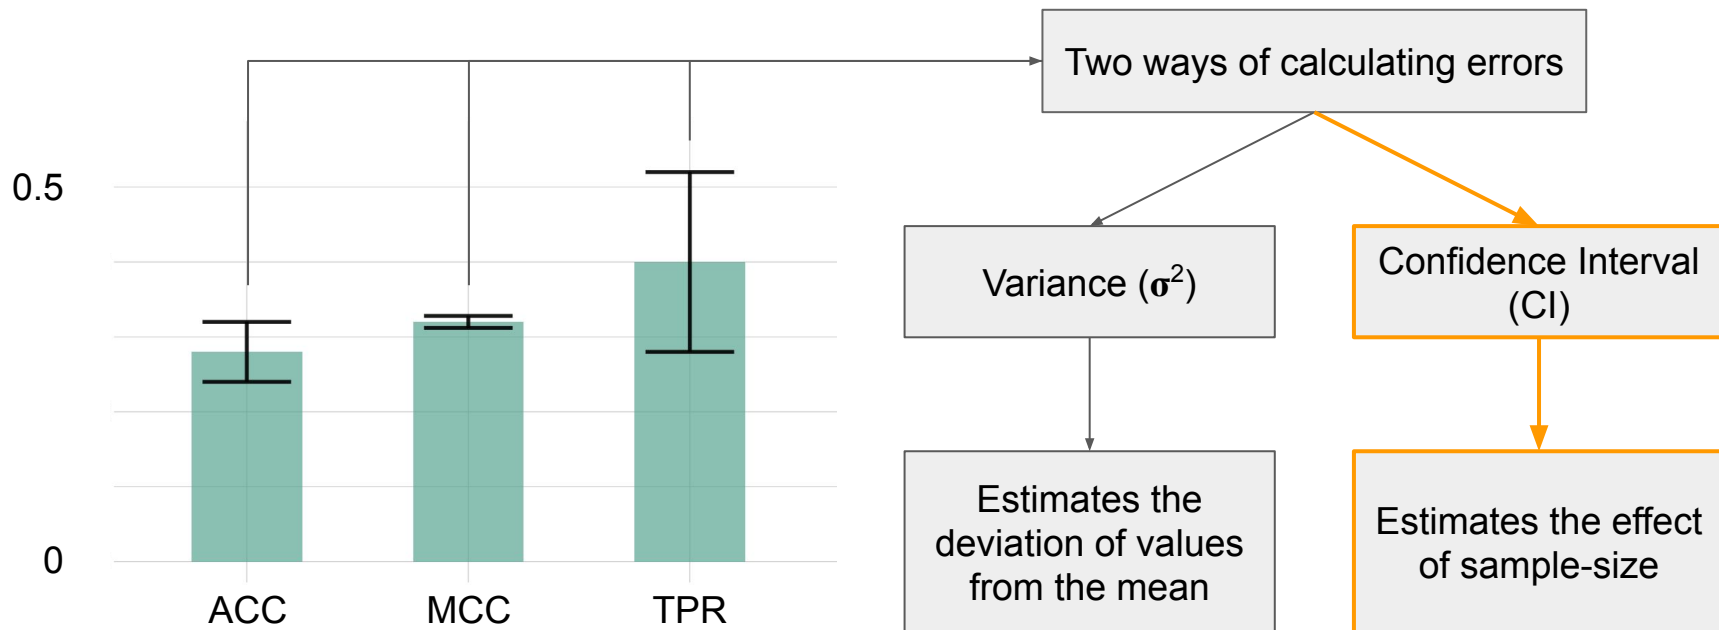
Dimension	Elements in dimension	Cumulative product	Order of magnitude	Execution time
Metrics	10	10	10^1	~1 second
Proteins	646	~6,500	10^3	~1 minutes
Thresholds	5	~32,500	10^4	~20 minutes

Code complexity

Dimension	Elements in dimension	Cumulative product	Order of magnitude	Execution time
Metrics	10	10	10^1	~1 second
Proteins	646	~6,500	10^3	~1 minutes
Thresholds	5	~32,500	10^4	~20 minutes
Predictors	30	~1,000,000	10^6	~1,5 hour

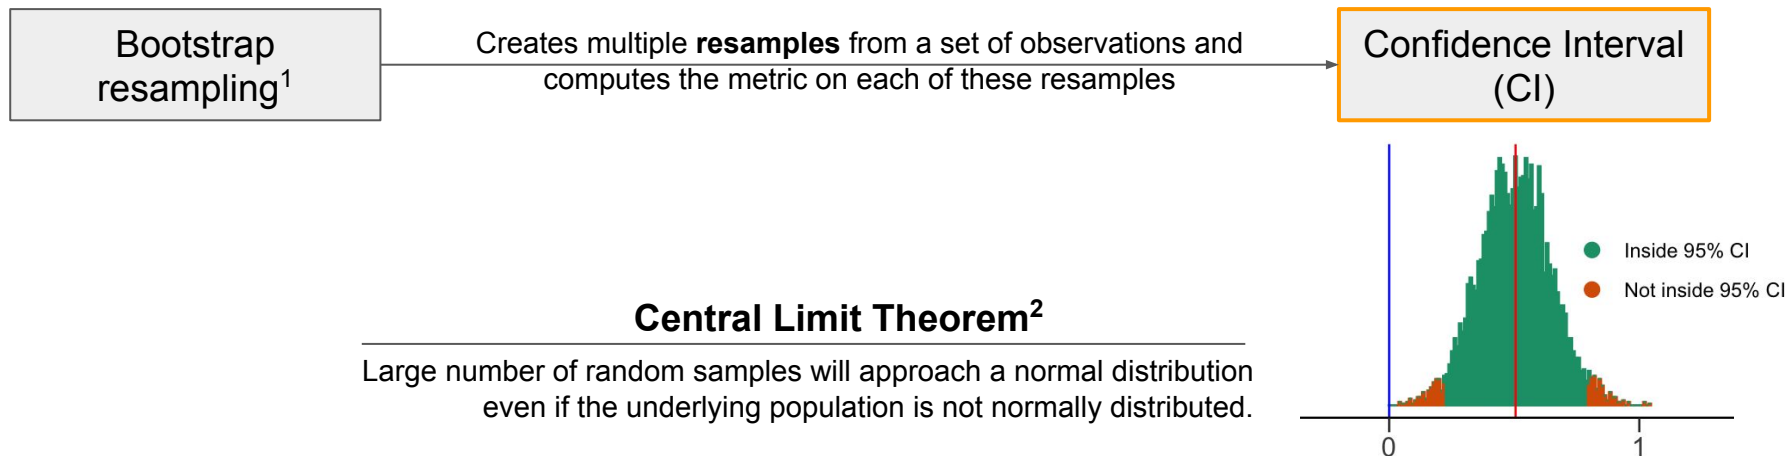
Confidence Intervals

- We want to evaluate the error associated to a metric



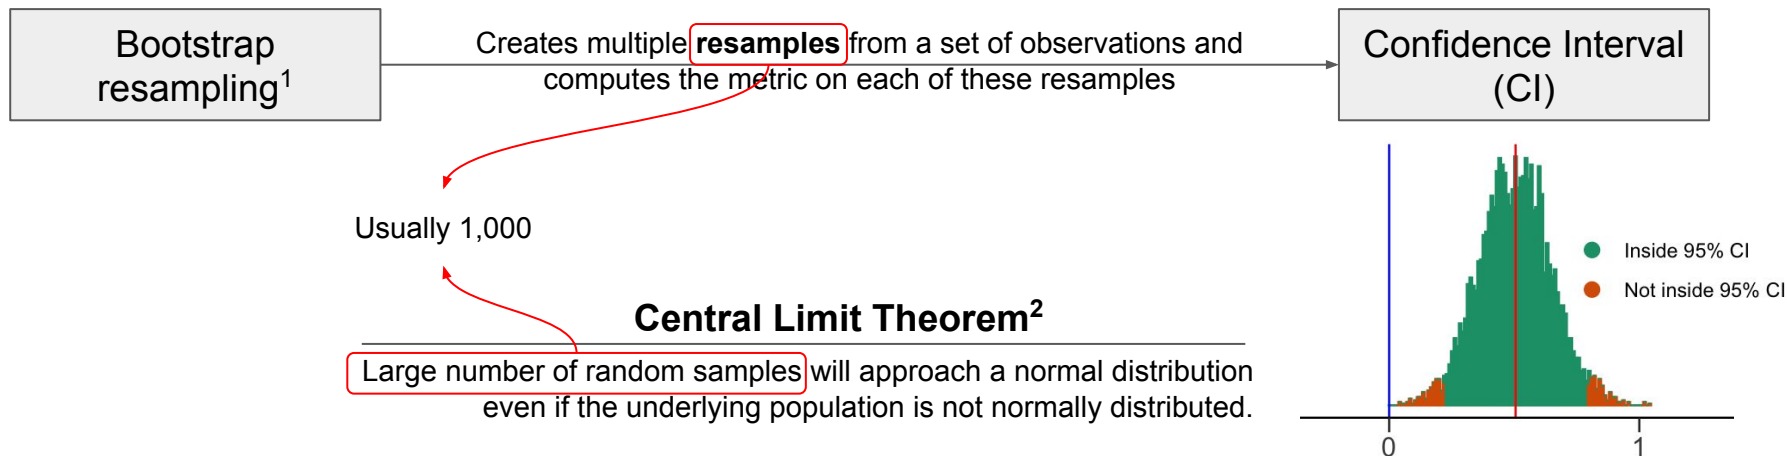
Confidence Intervals

- We want to evaluate the error associated to a metric



Confidence Intervals

- We want to evaluate the error associated to a metric



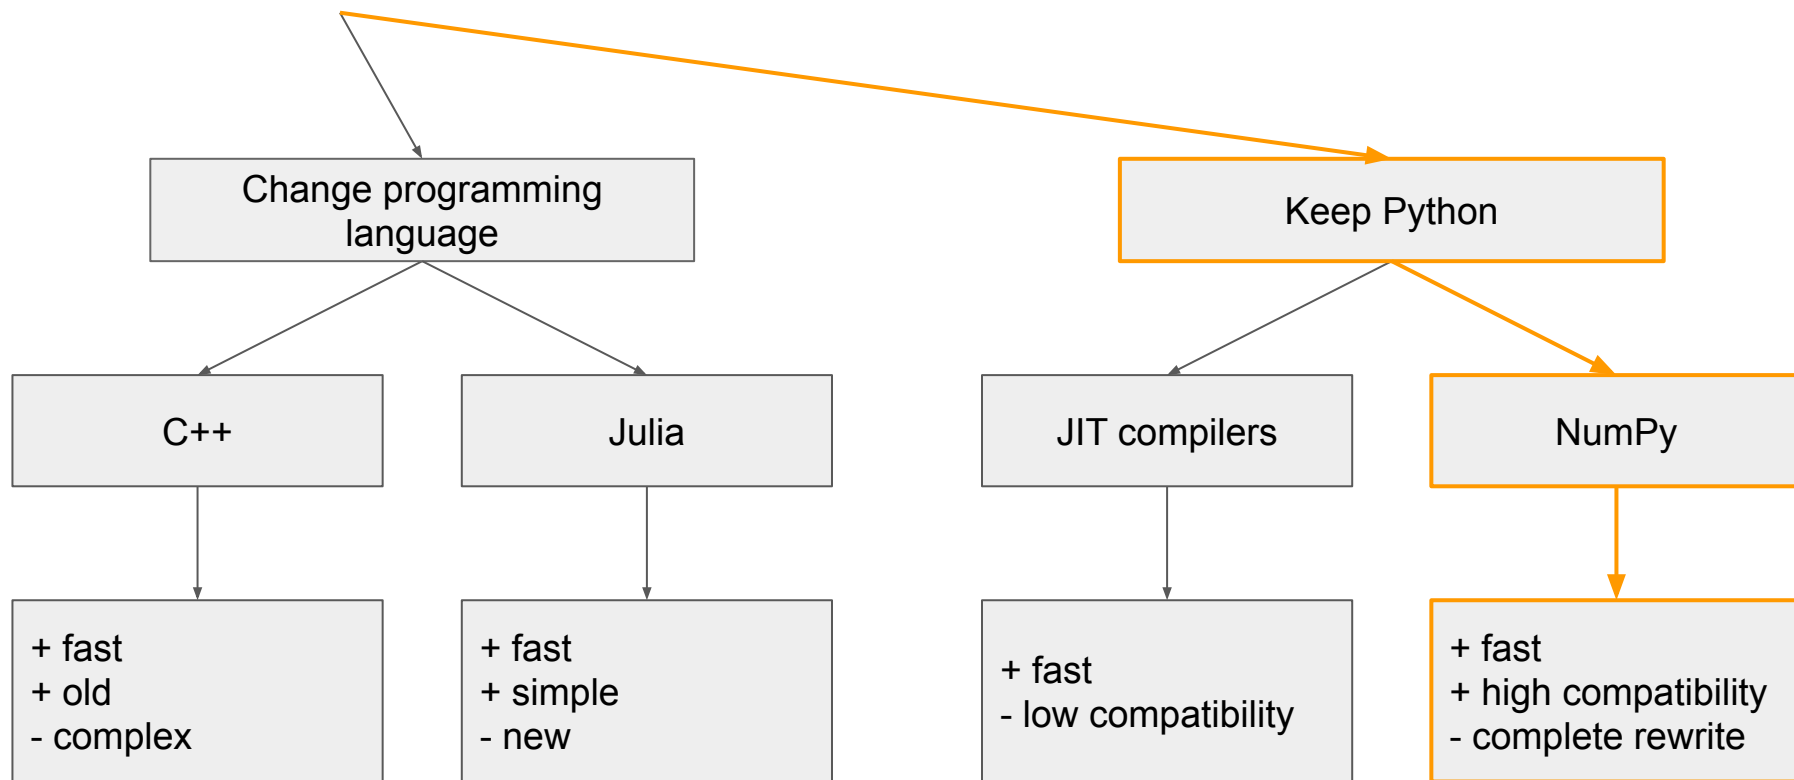
Code complexity

Dimension	Elements in dimension	Cumulative product	Order of magnitude	Execution time
Metrics	10	10	10^1	~1 second
Proteins	646	~6,500	10^3	~1 minutes
Thresholds	5	~32,500	10^4	~20 minutes
Predictors	30	~1,000,000	10^6	~1,5 hours
Bootstrap	1,000	~1,000,000,000	10^9	~15 hours

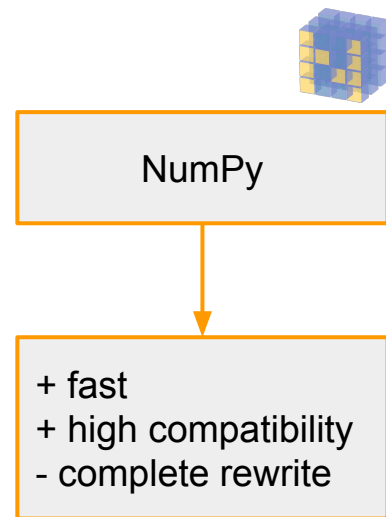

Solution

NumPy

Possible solutions



Vectorization



Vectorization

Calculates confusion matrix at all relevant thresholds

```
def binary_clf_curve(y_true, y_score):
    pos_label = 1.0
    y_true = (y_true == pos_label)


    desc_score_indices = np.argsort(y_score, kind="mergesort")[::-1]
    y_score = y_score[desc_score_indices]
    y_true = y_true[desc_score_indices]

    threshold_idx = np.r_[np.where(np.diff(y_score))[0], y_true.size - 1]

    tps = np.cumsum(y_true, dtype=np.float64)[threshold_idx]
    fps = 1 + threshold_idx - tps
    thr = y_score[threshold_idx]

    return fps, tps, thr
```


Vectorization



true

pred

Vectorization



```
array([0, 0, 0, 1, 1, 1, 0, 1, 1, 1])  
array([0.523, 0.957, 0.489, 0.244, 0.311, 0.093, 0.210, 0.065, 0.552, 0.556])
```


```
array([1, 9, 8, 0, 2, 4, 3, 6, 5, 7])
```

true

pred

Get the indexes of
pred sorted in
descending order

Vectorization

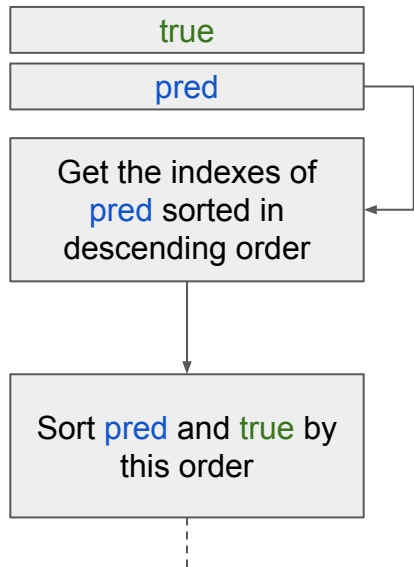


```
array([0, 0, 0, 1, 1, 1, 0, 1, 1, 1])  
array([0.523, 0.957, 0.489, 0.244, 0.311, 0.093, 0.210, 0.065, 0.552, 0.556])
```

```
array([1, 9, 8, 0, 2, 4, 3, 6, 5, 7])
```

All relevant thresholds

```
array([0.957, 0.556, 0.552, 0.523, 0.489, 0.311, 0.244, 0.210, 0.093, 0.065])  
array([0, 1, 1, 0, 0, 1, 1, 0, 1, 1])
```



All relevant thresholds

array([0.957, 0.556, 0.552, 0.523, 0.489, 0.311, 0.244, 0.210, 0.093, 0.065])

array([0, 1, 1, 0, 0, 1, 1, 0, 1, 1])

array([0, 1, 2, 2, 2, 3, 4, 4, 5, 6])

pred / thr

true

Sum true cumulatively

TPs

All relevant thresholds

array([0.957, 0.556, 0.552, 0.523, 0.489, 0.311, 0.244, 0.210, 0.093, 0.065])

array([0, 1, 1, 0, 0, 1, 1, 0, 1, 1])

array([0, 1, 2, 2, 2, 3, 4, 4, 5, 6])

array([1, 1, 1, 2, 3, 3, 3, 4, 4, 4])

pred / thr

true

Sum true cumulatively

TPs

Subtract TPs to the
count of elements

FPs

All positive classifications

array([0.957, 0.556, 0.552, 0.523, 0.489, 0.311, 0.244, 0.210, 0.093, 0.065])

array([0, 1, 2, 2, 2, 3, 4, 4, 5, 6])
array([1, 1, 1, 2, 3, 3, 3, 4, 4, 4])

array([3, 3, 3, 2, 1, 1, 1, 0, 0, 0])

All negatives

pred / thr

TPs

FPs


Subtract each FP
from the last FP

TNs

array([0.957, 0.556, 0.552, 0.523, 0.489, 0.311, 0.244, 0.210, 0.093, 0.065])

array([0, 1, 2, 2, 2, 3, 4, 4, 5, 6])
array([1, 1, 1, 2, 3, 3, 3, 4, 4, 4])

All positives



array([3, 3, 3, 2, 1, 1, 1, 0, 0, 0])

array([6, 5, 4, 4, 4, 3, 2, 2, 1, 0])

pred / thr

TPs

FPs

Subtract each FP
from the last FP

TNs

Subtract each TP
from the last TP

FNs

```
array([0, 1, 2, 2, 2, 3, 4, 4, 5, 6])
array([1, 1, 1, 2, 3, 3, 3, 4, 4, 4])
array([3, 3, 3, 2, 1, 1, 1, 0, 0, 0])
array([6, 5, 4, 4, 4, 3, 2, 2, 1, 0])
```

Confusion matrix at
this threshold

```
array([0.957, 0.556, 0.552, 0.523, 0.489, 0.311, 0.244, 0.210, 0.093, 0.065])
```

All relevant thresholds

TPs

FPS

TNs

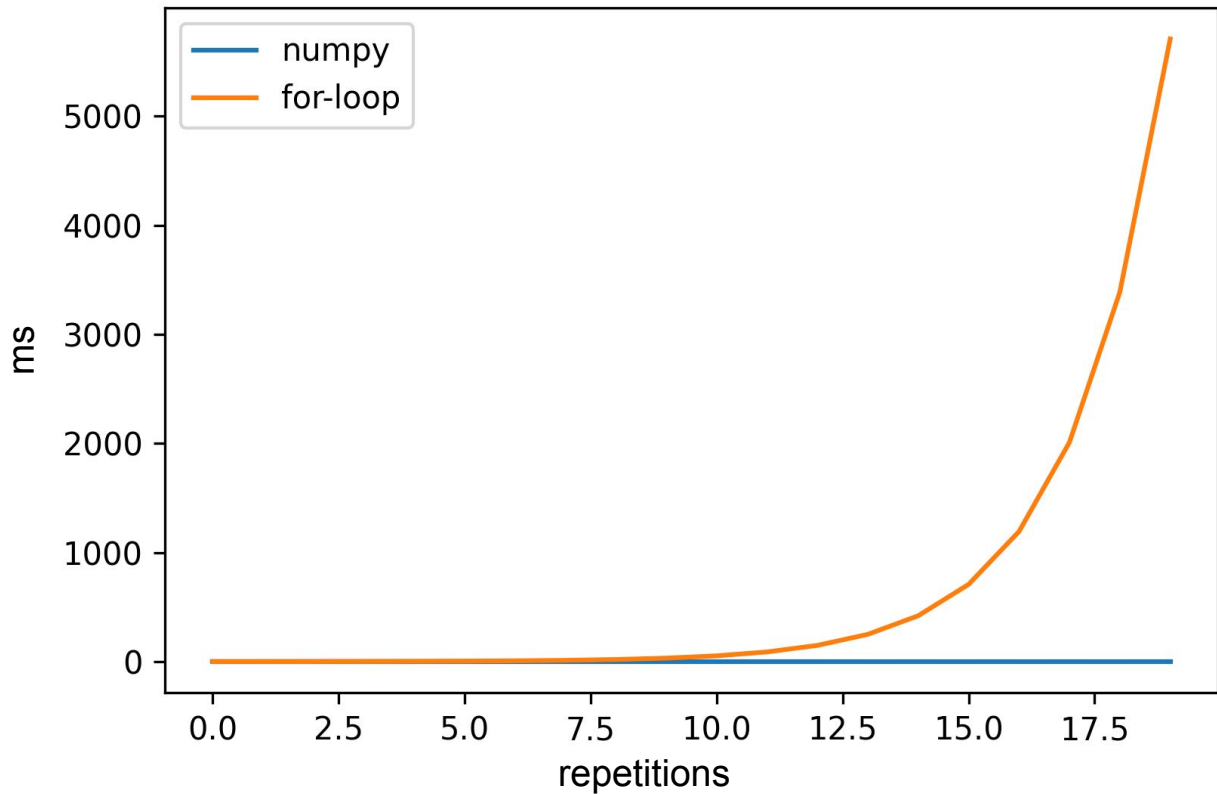
FNs

pred / thr

Results

Now it's fast

Efficiency



Curve of 20 points from linearly spaced x-values

Fitted on

Model from 5 points from exponentially spaced x-values

Code complexity

Dimension	Elements in dimension	Cumulative product	Order of magnitude	Execution time
Metrics	10 → 18	18	10^1	<1 second
Proteins	646	~11,500	10^4	<1 second
Thresholds	5 → 1000	~11,500,000	10^7	~1 minute
Predictors	30	~350,000,000	10^8	~1 minute
Bootstrap	1,000	~350,000,000,000	10^{11}	~10 minutes

Thank you